

## Abstract

# Parallelizing the Haplotyping Assembly Problem

Robert J. clucas \*

July 25, 2015

## 1 Introduction

It is commonly accepted that all humans share  $\sim 99\%$  of the same DNA, however, small variations cause human beings to have different physical traits. Single nucleotide polymorphisms (SNPs), which are variations of a single DNA base from one individual to another, are believed to be able to address genetic differences. For diploid organisms, which have pairs of chromosomes, a *haplotype* is a sequence of SNPs in each copy of a pair of chromosomes. A *genotype* describes the conflated data of the haplotypes on a pair of chromosomes. Haplotypes are believed to contain more generic information than genotypes [1], however, obtaining haplotypes correctly is a difficult problem, which has been studied in two main forms : haplotype inference and haplotype assembly (HA).

Haplotype inference uses the genotype of a set of individuals. The genotype data tells the status of each allele at a position, but does not distinguish which copy of the chromosome the allele came from. This negative aspects of this approach are that it cannot distinguish rare and novel SNPs [2], and there is no way of knowing if the inferred haplotype is completely correct.

Haplotype assembly uses fragments of sequences generated by sequencing technology to determine haplotypes. The fragments of a sequence come from the two copies of an individual's chromosome, the goal of the haplotype assembly problem is then to correctly determine two haplotypes, where each haplotype corresponds to one of the two copies of the chromosome. Table 1 in Appendix A shows an example input, where  $h$  and  $h'$  are the assembled haplotypes.

The haplotype assembly problem was proven to be NP hard [3]. The algorithms used to solve

the problem are thus computationally complex and until recently, there was no practical exact algorithm to solve the minimum error correction (MEC) formulation of the problem [4], However, recently an exact solution was proposed by [5] which is capable of solving the MEC formulation exactly, and can thus correctly infer all haplotypes from the fragment sequences. Due the NP hardness of the problem, the algorithm results in long run times - in the range of days for chromosomes with high errors rates. A parallel implementation of any of the proposed solutions could reduce the long run times, allowing useful haplotype information to be more quickly inferred from the available datasets, having positive effects in fields such as drug discovery, prediction of diseases, and variations in gene expressions, to name a few.

Parallel programming makes use of devices which have numerous cores, and uses these cores to execute a single instruction on multiple data (SIMD). The effectiveness of parallel programming is dependant on the nature of the problem, as per Amdahl's law. While using multiple processors can potentially provide large performance increases, in practise the performance improvements are difficult to achieve due to additional complexities which are introduced by parallelism. The main difficulties are the synchronization and communication between the multiple cores, and the management of memory between the host (normally a CPU) and the device (parallel capable hardware, a GPU for example). Threads are created to operate on some data in parallel, and a block of threads is allocated memory which is accessible only to the threads in the block, while all threads from all blocks have access to a global memory space, however, access to this memory space is slower, decreasing performance. Furthermore, race conditions, where multiple threads attempt to access data at the same location simultaneously, can cause undefined behaviour and hence inac-

---

\*School of Electrical & Information Engineering, University of the Witwatersrand, Private Bag 3, 2050, Johannesburg, South Africa

curate results.

With the increasing popularity of General-Purpose GPU (GPGPU) programming, however, API's like CUDA [6] and OpenCL [7] which provide access to GPUs through simple function calls in C and C++ program, have made parallel programming easier on GPUs. Intel have also made a move towards parallel programming, introducing the Xeon Phi, which, like a GPU, has numerous cores which can operate on data in parallel. The Xeon Phi is also programmed using an API [8], however, far more extensive use is made of compiler directives than by CUDA or OpenCL, with API functions being provided to gain access to parallel variables such as the thread index and number of threads which are running in parallel. Using these APIs, it is possible to provide parallel implementations of algorithms which have significantly shortened run times.

The contribution of this paper is to consider the proposed algorithms for solving the HA problem, then to choose one for which parallelization can potentially provide performance improvements, and then to identify components of the selected algorithm which would be suited for parallelism as well as proposing possible parallel implementation of identified parallelizable components. Additionally, this paper details the specification of the project and breakdown of work such that the project will be completed by the deadline.

The remainder of this paper is structured as follows. Section 2 provides background information related to the HA problem, and some of the proposed algorithms for solving it, the MEC formulation, and the branch and bound algorithm. Section 3 describes the specifics of the project. Section 5 describes possible parallelization of the preprocessing of the input data. Section 6 describes the parallelization of the branch and bound algorithm, which is the chosen algorithm for solving the problem. Section 7 concludes.

## 2 Background

### 2.1 Haplotype Assembly Problem

This subsection will provide a brief overview of the haplotype assembly problem, and define the notation used through the rest of the paper. The input to the problem is a set of reads from a

given genome sequence, where each read contains fragments from each of the two chromosomes which make up the genome sequence. The characters of a read consist of elements from a *ternary string*, where a ternary string has characters from the set  $\{0, 1, -\}$ . A value of 0 refers to the major allele at a site, a value of 1 to the minor allele, and a value of - to the lack of a read at the site, and is referred to as a *gap*. These reads are then combined to form a matrix,  $M$ , where each row of the matrix corresponds to a read (Table 1 in Appendix A provides an example).

Each column of the matrix is known as an SNP site. At each site, the data could be accurate, missing, or have error. The goal of the haplotype assembly problem is to determine a haplotype,  $H = \{h, h'\}$  from the matrix. The following terminology will be used to refer to properties of the matrix and the fragments.

For the input matrix,  $M$ , the number of fragments is denoted by  $m$ , which is the number of rows in  $M$ . The number of SNP sites is denoted by  $n$ , which is the number of columns in  $M$ , while the  $j^{th}$  site of the  $i^{th}$  fragment is given by  $f_{ij}$ . Two fragments are said to conflict if the following conditions are true:

- $f_{ik} \neq f_{jk}$  and  $f_{ik} \neq '-'$  and  $f_{jk} \neq '-'$

Essentially this means that for two fragments  $i$  and  $j$ , if at an SNP site  $k$ , the reads do not have error and are not gaps, the reads have different values (fragment  $i$  has a value 0 at site  $k$ , while fragment  $j$  has a value 1 at site  $k$ , or vice versa).

Following the notion of a conflict, the *distance* between two fragments is denoted by  $d(f_i, f_j)$  is the total number of positions for which the two fragments  $f_i$  and  $f_j$  conflict. Furthermore, to understand some of the problem formulations from the fragment data, it is useful to define a *conflict graph*  $G = \{V, E\}$  [9], where  $V$  corresponds to a fragment, and  $E$  corresponds to an edge between two fragments if they conflict. If the input matrix contains no errors, then none of the fragments from the same chromosome will conflict and  $G$  will be bipartate. However, if there are errors (as is usually the case) in  $M$ ,  $G$  will not be bipartate. The haplotype assembly problem then requires the correction of  $G$  from a non-bipartate graph to a bipartate graph, from which two haplotypes can be assembled. There are numerous methods for solving the problem:

- **Minimum Fragment Removal (MFR):**  
This involves removing the least number of

fragments from the input data such that the resultant graph  $G$  is bipartate. It is shown in [9] that this can be solved in polynomial time.

- **Minimum Edge Removal (MER) [10]:** This method was recent proposed, and requires determining the minimum number of edges to remove such that removal of the edges results in  $G$  being bipartate.
- **Longest Haplotype Reconstruction (LHR):** This requires finding a set of fragments which, when they are removed from  $M$ , result in  $G$  being bipartate and the length of the resultant haplotypes being maximized [11].
- **Minimum Error Correction (MEC):** This method involves correcting the minimum number of elements (sites for all fragments) in the matrix  $M$  which allows the graph  $G$  to be bipartate. Although being the most complex method, it is the most widely used method as it provides the highest accuracy. Only recently has an exact algorithm for the MEC formulation been proposed which can provide an exact solution for all cases.

Due to the accuracy and more extensive previous work, the MEC formulation of the HA problem was chosen for a parallel implementation, hence will be the focus for the remainder of the paper.

## 2.2 Minimum Error Correction Formulations

Much research has been done on solving the MEC formulation of the HA problem. The first exact algorithm for solving the problem was a branch and bound algorithm proposed by [12]. The algorithm creates a tree which covers the search space of all possible corrections. The tree is then traversed to find the best solution. They use an upper bound which allows branches to be pruned when a better solution than the current best cannot result from further exploration of the branch, resulting in faster run times. However, this exact method has time complexity of  $O(2^m)$ , where  $m$  is the number of fragments in the input data, and hence cannot determine solutions for large problem sizes within a feasible amount of time. They also provide a heuristic method which uses a *genetic algorithm* to improve the computational time, with run times

up to three orders of magnitude faster than the branch and bound implementation. While the heuristic method gives very similar results to the branch and bound implementation, it is slightly less accurate.

A dynamic programming solution was proposed by [13] which addresses the run time problem for large input sizes. The algorithm has time complexity of  $O(mk3^k + m\log m + mk)$ , where  $k$  is the maximum number of SNP sites a fragment covers. In practice  $k$  is usually small, and results were shown small  $k$  (less than 100). The proposed dynamic programming solution had significant run time improvements over the solution proposed by [12]. However, for larger  $k$  values, the algorithm cannot solve the MEC formulation of the HA problem within a feasible amount of time.

More recently, [5] proposed an exact algorithm for solving the MEC problem. The proposed algorithm is the currently the only algorithm which can solve the HA problem for both the *heterozygous* case (the alleles at an SNP site are different assumed to be different) and the *general* case (the alleles at an SNP site may be the same due to error). Most other work assumes that the input fragment data is heterozygous, which, while true for most of the SNP sites in the input matrix, is often false for a small number of the SNP sites, which is the benefit of the algorithm proposed by [5]. Their exact algorithm removes unnecessary data from the input matrix, and partitions the matrix into smaller sub-matrices (or unsplittable blocks to use their terminology, Appendix A provides an example) which can then be solved in parallel. The problem is then formulated as an *integer linear programming* (ILP) problem and solved as a minimization problem. Their implementation was tested using an Intel i7-3960X CPU, and the HuRef dataset required 12 days to solve for the general case, and 31 hours for the all-heterozygous case. Heuristics methods are also proposed which speed up computation time by up to 15 times. However, the results are only shown for smaller input sizes, and the heuristic methods does not always determine the optimal solution. Most importantly, the solutions for the general case show lower MEC scores, meaning that the all-heterozygous assumption is not always valid. This algorithm for the MEC formulation was chosen as it achieved optimal solutions for the all heterozygous and general cases. Their exact formulation of the problem as an

ILP problem for the all-heterozygous case is now given.

### 2.2.1 All-heterozygous Algorithm by Chen

The Hamming distance between two fragments  $j$  and  $k$ ,  $d = (f_i, f_j)$ , is used by [5] to determine the MEC score of a solution  $H = \{h, h'\}$ , and is number of SNP sites at which the fragments conflict. Using the Hamming distance between all the fragments and the solutions, the MEC score is given by  $d(f_j, f_k)$

$$\text{MEC score} = \sum_{i=1}^m \min\{d(f_i, h), d(f_i, h')\} \quad (1)$$

Where  $m$  is again the number of rows in the input matrix  $M$ . The MEC score is *optimal* if it is the minimum possible score. Some assumptions are made for the input matrix as per [5], namely

- No row of the input matrix  $M$  is useless - at least one entry in the row is a 1 or 0
- No column of  $M$  is monotone - the column must have at least one 0 and one 1
- No column of  $M$  contains more 1's than 0's - if this is not the case the values are all flipped, which does not change the solution or the MEC score

The input matrix first undergoes pre-processing to reduce the size of the input and to break the input into multiple inputs which can be solved in parallel. It must be noted that the pre-processing does not affect the solution in any way. The pre-processing functions are:

- **Block decomposition:** This process takes the input matrix and splits the input into smaller, unsplitable blocks, which are disjoint and can be solved independently.
- **Singleton removal:** A singleton is a row for which the start and end positions of the fragment are the same - i.e the fragment has only one element. Since singletons do not affect the MEC score, they can be removed. This is done for all rows in all the smaller, unsplitable blocks.
- **Duplicate removal:** Rows and columns which are the same are merged into a single row or column, and the multiplicity of the row or column is recorded. This is done for

all rows and all columns in all the smaller, unsplitable blocks.

From the reduced blocks the ILP problem is formulated for the all-heterozygous and the general case. The all heterozygous case formulation will be shown, the general case formulation is given shown in CITE.

**2.2.1.1 ILP Formulation** For integers  $p, q \in \mathbb{Z}$ , where  $1 \leq p \leq m$ , and  $1 \leq q \leq n$ ,  $p$  is the index of the fragment (or row) in the input matrix  $M$ , and  $q$  is the index of SNP (or column) in the input matrix. The multiplicity of the  $j^{th}$  column of  $M$  is denoted by  $c_j$ , while the multiplicity of the  $i^{th}$  row of  $M$  is denoted by  $c_i$ , an example is provided in Appendix B. The following binary variables are introduced,  $y_i$ , which has a value of 1 if  $d(f_i, h) < d(f_i, h')$ , otherwise has a value of 0 (when  $d(f_i, h) < d(f_i, h')$ ). Informally, if fragment  $i$  is part of  $h$  then  $y_i$  has a value of 1, otherwise it has a value of 0, and  $x_j$ , which has a value of 1 if the  $j^{th}$  bit of  $h$  is 1, otherwise has a value of 0 of the  $j^{th}$  bit of  $h$  is 0. Lastly  $J_{i,0}$  ( $J_{i,1}$ ) are the sets of integers  $j \in \{1, 2, \dots, q\}$  for which the  $i^{th}$  value in column  $j$  is a 0 (1), an example is given in Appendix B.

Using the above mentioned variables, an integer programming formulation is possible, however, a non-linear term in the form of  $y_i x_j$  arises, which cannot be solved using ILP techniques. To overcome this problem, [5] defines a variable  $t_{i,j}$  for  $y_i x_j$  and impose constraints on the variables which ensure linearity (the constraints are shown in the final formulation of the problem below). Using these variables, the final ILP formulation of the HA problem for the all heterozygous case is

$$\begin{aligned} \text{Minimize} \quad & \sum_{i=1}^p w_i \sum_{j \in J_{i,0}} c_j (1 - x_j - y_i + 2t_{i,j}) \\ & + \sum_{i=1}^p w_i \sum_{j \in J_{i,1}} c_j (y_i + x_j - 2t_{i,j}) \end{aligned}$$

$$\begin{aligned} \text{Subject to} \quad & \forall_{1 \leq i \leq p} \quad y_i \in \{0, 1\} \\ & \forall_{1 \leq j \leq q} \quad x_j \in \{0, 1\} \\ & \forall_{1 \leq i \leq p} \quad \forall_{1 \leq j \leq q} \quad t_{i,j} \in \{0, 1\} \\ & t_{i,j} \leq y_i \\ & t_{i,j} \leq x_j \\ & t_{i,j} \geq y_i + x_j + 1 \end{aligned}$$

## 2.3 Branch and Bound

The branch and bound algorithm divides the search space into sub spaces using a branching operator, and each sub space is explored for the optimal solution. Each of these sub spaces is a branch on a tree, a bounding operator is used to determine the best possible solution the branch can provide. A pruning operator is used to remove branches which cannot provide a solution which is better than the current best, based on the result of the bounding operator. The branch and bound algorithm can be parallelized in multiple ways [14]. Either specific, computationally difficult functions can be accelerated by a parallel implementation (for example matrix inversion), or the entire tree can be searched in parallel, or a combination of both can be employed. Using tree based approaches will require communication between the processes searching a branch, as the solution of each branch will need to be compared to the solutions of the other branches to determine the globally optimal solution. This can become a bottleneck for both CPU and GPU implementations if the work is not divided correctly as processes may have to wait for other processes. There have been numerous parallel branch and bound algorithms proposed for both the CPU and GPU which deal with these problems in different ways.

### 2.3.1 CPU Implementations

The ALPS framework [15] is written in C++ and provides a parallel CPU implementation of the branch and bound algorithm. The framework is tested on the knapsack problem [16], which is an ILP problem and is known to be NP hard. The speedup achieved is near linear in the number of nodes when the number of nodes is small, but diminishes as more nodes are added due to the amount of communication between the nodes. A speedup of 8 times is achieved for 8 nodes, and 26 times for 32 nodes. It is likely that for the HA problem, the number of nodes could be extremely large in which case this method will be infeasible.

The MALLBA framework [17] is also written in C++ and provides a parallel branch and bound algorithm. It labels nodes as *master* or *slave* nodes, which defines the type of work that the nodes do. The framework also provides heuristic methods for solving the problems more efficiently, however, the accuracy of the results

is lowered, which while acceptable for many applications, is not acceptable for the HA problem. The performance results are similar to the ALPS framework, where near linear speed up is achieved for a small number of additional nodes, but levels off for larger numbers of nodes.

### 2.3.2 GPU Implementations

A heterogeneous CPU-GPU implementation was proposed by [18] and was applied to the knapsack problem. Due to the overhead of transferring data from the CPU to the GPU before computation, the model only uses GPUs when the tree has a large number of nodes ( $> 5000$ ), otherwise the CPUs are used. The tree is built using a *breadth first* strategy to favour the parallel nature of the GPUs. Both the branching and bounding steps can be performed on the CPU or GPU. If the number of nodes is sufficiently large such that full occupancy of the GPU is ensured, then for each iteration of the algorithm the GPU does the branching and bounding on a list of nodes, eliminates nodes for which an optimal solution cannot be found, and returns the list back to the CPU. This process continues until convergence. This regular communication between the CPU and GPU is expensive, which lessens the speedup achieved by the implementation. However, a speedup of up to 9.27 times was achieved for larger problem sizes, validating the feasibility of the branch and bound algorithm for parallel implementation. It must be noted that the algorithm was developed for Nvidia's Fermi architecture, which does not support dynamic parallelism [19], thus requires numerous CPU-GPU communication. The newer Kepler architecture, however, does support dynamic parallelism and could minimise this communication.

In [20, 21], algorithms are proposed which target the bounding operation for parallelism, as well as dealing with the problem of thread divergence when using GPUs for branch and bound, which comes from the position of the nodes in the tree, and the need to for branches to communicate with other branches to compare solutions. Depending on the problem to which branch and bound is applied, the tree structure can be irregular which makes parallelizing the tree search difficult for GPUs not supporting recursion, as was the case when the algorithm was proposed – hence the focus on only the bounding operation. The algorithm was applied to the

Flow-Shop scheduling problem, for which it is shown that  $\sim 98\%$  of the computation is spent on the bounding operation. Speed ups of up to 100 times were achieved by the GPU implementation over a multi-threaded CPU implementation. However, this kind of speedup will only be seen in problems where the calculation of the bounds is the bottleneck. Nevertheless, even if significantly less time is spent calculating the bounds, the speed up should still be considerable.

The algorithms of [20, 21] are extended by [22] to include not only parallelization of the bounding operator, but also the branching and pruning operators. This implementation uses the GPU for almost all the searching of the subspaces. However, the CPU is still used for the comparison of the solutions between each iteration as each GPU thread can only determine the solution of its first child nodes, rather than all child nodes until the leaves of the tree. Again this is due to the limitations of the Fermi architecture not supporting recursion. Despite this hardware limitation, the algorithm is still able to achieve a speed up of up to 166 times over a multi-threaded CPU implementation, for large problem sizes.

## 3 Project Description

Based on the related work reviewed in Section 2, the specification of the project and its purpose, the assumptions, and constraints and requirements, are given in the subsections to follow.

### 3.1 Project Specification

The aim of the project is to implement the HA problem in parallel on 3 different hardware configurations and then to compare the performance of the configurations. The chosen algorithm for parallel implementation is the MEC formulation proposed by [5] since the algorithm can solve the all heterozygous and general case optimally. The 3 configurations are:

- CPU based cluster
- Intel Xeon Phi Coprocessor based cluster
- Nvidia GPU based cluster

### 3.2 Hardware

The following hardware will be required for the project:

- A cluster with multi-core CPUs
- A cluster with at least one Intel Xeon Phi Coprocessor
- A cluster with at least one Nvidia GPU with compute capability  $\geq 3.5$

### 3.3 Software

The following software will be used for the project:

- C and C++, as these are the languages used for programming Intel Xeon Phi's and Nvidia GPU's
- The Nvidia CUDA API for programming the Nvidia GPUs
- The OpenMP API for the CPU and Intel Xeon Phi implementations
- The Nvidia nvcc compiler for compiling the GPU code to run on the Nvidia GPU
- The Intel compiler suite for compiling the CPU and Intel Xeon Phi implementations

### 3.4 Assumptions

The following assumptions are made for the project:

- The hardware detailed in Section 3.2 is available
- The software detailed in Section 3.3 is available
- The datasets for the SNP inputs will all be available
- Solutions for the datasets are available to verify the correctness of the implementations

### 3.5 Constraints

The following constraints are placed on the project:

- The available budget is limited and thus a limited number of hardware devices will be available
- All software to be used must be free
- The project must be completed by Robert Clucas and Sasha Naidoo

### 3.6 Success Criteria

Due to the difficulty of the problem, the following criteria are defined which must be met for the project to be considered a success:

---

**Algorithm 1** Procedure for solving the MEC HA problem using ILP

---

**Step 1:** Perform block decomposition on columns in  $M$  to obtain smaller unsplittable blocks  $C_k$  (if general case, check if column is intrinsically heterozygous)

**Step 2:** Remove singleton rows from all  $C_k$  unsplittable blocks in parallel (if general case, check if singleton row starts and ends on intrinsically heterozygous columns)

**Step 3:** Remove duplicate columns from all  $C_k$  unsplittable blocks

**Step 4:** Solve ILP formulation using branch and bound algorithm on all  $C_k$  unsplittable blocks

---

- All three implementations are working by the project demonstration date
- All implementations can at least solve the all heterozygous case correctly, with optimal MEC scores
- At least one implementation provides a speed up over the results detailed in [5]

## 4 Proposed Implementation

Using the information provided by Section 2.2.1 and Section 2.2.1.1, the general procedure for solving the MEC HA problem using ILP can be defined as

## 5 Parallel Pre-processing for MEC HA

Observing Algorithm 1, as well as referring to [5] for the exact definitions of the operations, there are numerous areas to which parallelism can be applied.

### 5.1 Determining Intrinsically Heterozygous Columns

Determining intrinsically heterozygous columns in  $M$  requires comparing counting the number of rows at each of the  $j$  columns which are 0 and which are 1, and then comparing the value. A parallel implementation would perform this operation for all  $n$  columns at once if the GPU was used. However, for small  $n$  the overhead of transferring the data to the GPU for computation would result in worse performance. In this case, a better approach would be to use the vectorized operations

available to modern CPUs to perform the operations on multiple columns at a time, which could result in up to 8 times speedup depending on the CPU.

### 5.2 Singleton Removal

Finding singletons is a simple task as it only requires determining if there is more than one element in a row. Again, vectorized operations could be used, however, this time one the GPU, which could use  $n$  threads for each row. Each thread checks if its corresponding entry in the row is present, and if it is, increments a counter in shared memory. If the value of the counter is  $> 1$ , then the row is not singular. This would have time complexity of  $O(1)$ , compared to  $O(n)$  using the implementation of [5].

### 5.3 Duplicate Removal

For the implementation used by [5], each row needs to be compared with each other row to determine if the rows are identical, and the same procedure is required for the columns. This can be done with time complexity of  $O(L \log k)$ , where  $L$  is the length of the reads and  $k$  is the number of reads. There are numerous ways in which this operation can be parallelized.

#### 5.3.1 Simultaneous column and row search

This would use the GPU to create threads that perform the column and row comparisons at the same time. Furthermore, multiple columns and rows can be done at the same time. For example, the  $i^{th}$  and  $z^{th}$  rows can be compared to all the other rows at the same time. This would allow a speed up on the order of

$$\text{Speed up} = N * \frac{\text{GPU frequency}}{\text{CPU frequency}} \quad (2)$$

where  $N$  is the number of threads that are used simultaneously.

#### 5.3.2 Simultaneous row and column comparison

This implementation would compare the  $i^{th}$  row or column to multiple other rows or columns at the same time (for example to the  $\{a^{th}, b^{th}, \dots, j^{th}, k^{th}, \dots, z^{th}\}$  rows or column at the



same time). This would be much faster provided the parallel device has a sufficient number of threads. For very large input sizes there will be a serial aspect as there will not be enough threads. Nevertheless, for each of the iterations, the time complexity would be  $O(L + K)$ , assuming there were  $2 \cdot L \cdot K$  threads available at each iteration, where  $L$  is again the length of the reads and  $K$  is the number of reads (this is because each row would spawn  $L \cdot K$  threads to cover the whole matrix on each iteration, and each column would spawn the same).

## 6 Parallel Branch and Bound for MEC HA

Considering the related work presented in Section 2.3, a similar approach to those presented in [20–22] will be taken, however, an attempt will be made to limit the number of transactions between the CPU and GPU, preferring rather to perform as much calculation on the GPU as possible. Furthermore, the implementation will not attempt to balance the tree since the parallel devices which will be used are capable of recursion, allowing branches to be explored until either an optimal or infeasible solution is reached. Simply splitting the problem space into many sub problems will not be sufficient for large performance increases. The hard blocks of the HuRef dataset took up to 12 hours on a multi-core CPU [5]. Operator level parallelism can potentially improve this.

### 6.1 Node Selection

Once a branch has been searched and cannot be explored further because it’s solution is either optimal or cannot be better than the current lower bound, the group of threads must be re-assigned to search a new branch. A depth first strategy will be applied since the parallel hardware supports recursion and hence can handle an unstructured tree.

### 6.2 Branching Operator

The branching operator is a kernel function that is called from any of the parallel compute devices. A block of threads is assigned a branch from the selection operator. The size of the block of threads can be determined statically from the depth of the branch, which will ensure

that the number of threads is carefully managed and that additional thread blocks will only be assigned if they are available. While this method ensures correct operation, it is sub optimal in terms of performance.

Alternatively, the number of threads can be determined dynamically, where the search from the root of the branch is started from a single thread, then the subsequent child nodes are searched by new threads created from the parent nodes. This process continues until an optimal or infeasible solution is found by one of the child threads, at which point the root thread is returned to before getting another branch to search. This method will always use exactly the correct number of threads, however, if the recursion becomes too deep across many branches, resources management could become a problem and ensuring the availability of resources could detract from the performance.

### 6.3 Bounding Operator

## 7 Conclusion

## References

- [1] J. Stephens. “Haplotype Variation and Linkage Disequilibrium in 313 Human Genes.” *Science*, vol. 293, no. 5529, pp. 489–493, Jul. 2001. [Online]. Available at <http://dx.doi.org/10.1126/science.1059431> [Accessed 27 June 2015].
- [2] D. He, A. Choi, K. Pipatsrisawat, A. Darwiche, and E. Eskin. “Optimal algorithms for haplotype assembly from whole-genome sequence data.” vol. 26, no. 12, pp. i183–i190, 2010.
- [3] R. Lippert, R. Schwartz, G. Lancia, and S. Istrail. “Algorithmic strategies for the single nucleotide polymorphism haplotype assembly problem.” vol. 3, no. 1, pp. 23–31, 2002.
- [4] P. Bonizzoni, G. Della Vedova, R. Dondi, and J. Li. “The Haplotyping problem: An overview of computational models and solutions.” *Journal of Computer Science and Technology*, vol. 18, no. 6, pp. 675–688, 2003. [Online]. Available at <http://>

- [//dx.doi.org/10.1007/BF02945456](http://dx.doi.org/10.1007/BF02945456) [Accessed 27 June 2015].
- [5] Z.-Z. Chen, F. Deng, and L. Wang. “Exact algorithms for haplotype assembly from whole-genome sequence data.” vol. 29, no. 16, pp. 1938–1945, 2013.
  - [6] Nvidia. *CUDA C PROGRAMMING GUIDE. Version 7.0*. 2015. [Online] Available at <http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#axzz3gRwjcm3m> [Accessed 20 July 2015].
  - [7] Khronos. *The OpenCL Specification. Version 2.1*. 2015. [Online] Available at <https://www.khronos.org/registry/cl/> [Accessed 20 July 2015].
  - [8] Intel. *Intel Xeon Phi Coprocessor DEVELOPER’s QUICKSTART GUIDE. Version 1.8*. 2013.
  - [9] G. Lancia, V. Bafna, S. Istrail, R. Lippert, and R. Schwartz. “SNPs Problems, Complexity, and Algorithms.” In F. auf der Heide, editor, *Algorithms ESA 2001*, vol. 2161 of *Lecture Notes in Computer Science*, pp. 182–193. Springer Berlin Heidelberg, 2001. [Online] Available at [http://dx.doi.org/10.1007/3-540-44676-1\\_15](http://dx.doi.org/10.1007/3-540-44676-1_15) [Accessed 16 July 2015].
  - [10] D. Aguiar and S. Istrail. “HAPCOMPASS: A fast cycle basis algorithm for accurate haplotype assembly of sequence data.” *Journal of Computational Biology*, vol. 19, no. 6, pp. 577–590, June 2012. [Online] Available at [http://www.brown.edu/Research/Istrail\\_Lab/papers/hapcompass\\_jcb\\_draft.pdf](http://www.brown.edu/Research/Istrail_Lab/papers/hapcompass_jcb_draft.pdf) [Accessed 16 July 2015].
  - [11] R. Schwartz. “Theory and Algorithms for the Haplotype Assembly Problem.” *Commun. Inf. Syst.*, vol. 10, no. 1, pp. 23–38, 2010. [Online] Available at <http://projecteuclid.org/euclid.cis/1268143371> [Accessed 16 July 2015].
  - [12] R.-S. Wang, L.-Y. Wu, Z.-P. Li, and X.-S. Zhang. “Haplotype reconstruction from SNP fragments by minimum error correction.” vol. 21, no. 10, pp. 2456–2462, 2005.
  - [13] M. Xie, J. Wang, and J. Chen. “A Practical Exact Algorithm for the Individual Haplotyping Problem MEC.” In *BioMedical Engineering and Informatics, 2008. BMEI 2008. International Conference on*, vol. 1, pp. 72–76. May 2008.
  - [14] T. G. Crainic, B. Le Cun, and C. Roucairol. *Parallel Branch-and-Bound Algorithms*, pp. 1–28. John Wiley & Sons, Inc., 2006. [Online] Available at <http://dx.doi.org/10.1002/9780470053928.ch1> [Accessed 17 July 2015].
  - [15] Y. Xu, T. Ralphs, L. Ladnyi, and M. Saltzman. “Alps: A Framework for Implementing Parallel Tree Search Algorithms.” In B. Golden, S. Raghavan, and E. Wasil, editors, *The Next Wave in Computing, Optimization, and Decision Technologies*, vol. 29 of *Operations Research/Computer Science Interfaces Series*, pp. 319–334. Springer US, 2005. URL [http://dx.doi.org/10.1007/0-387-23529-9\\_21](http://dx.doi.org/10.1007/0-387-23529-9_21). [Online] Available at [http://dx.doi.org/10.1007/0-387-23529-9\\_21](http://dx.doi.org/10.1007/0-387-23529-9_21) [Accessed 17 July 2015].
  - [16] M. Kedia. “Dynamic Programming.”, October 2005. [Online] Available at <http://www.cs.cmu.edu/afs/cs/academic/class/15854-f05/www/scribe/lec10.pdf> [Accessed 27 June 2015].
  - [17] E. Alba, F. Almeida, M. Blesa, J. Cabeza, C. Cotta, M. Daz, I. Dorta, J. Gabarr, C. Len, J. Luna, L. Moreno, C. Pablos, J. Petit, A. Rojas, and F. Xhafa. In B. Monien and R. Feldmann, editors, *EuroPar 2002 Parallel Processing*, vol. 2400 of *Lecture Notes in Computer Science*, pp. 927–932. Springer Berlin Heidelberg, 2002. [Online] Available at [http://dx.doi.org/10.1007/3-540-45706-2\\_132](http://dx.doi.org/10.1007/3-540-45706-2_132) [Accessed 17 July 2015].
  - [18] A. Boukedjar, M. Lalami, and D. El-Baz. “Parallel Branch and Bound on a CPU-GPU System.” In *Parallel, Distributed and Network-Based Processing (PDP), 2012 20th Euromicro International Conference on*, pp. 392–398. Feb.
  - [19] Nvidia. “CUDA DYNAMIC PARALLELISM PROGRAMMING GUIDE.”,

August 2012. [Online] Available at <https://www.clear.rice.edu/comp422/resources/cuda/html/cuda-dynamic-parallelism/index.html> [Accessed 17 July 2015].

- [20] N. Melab, I. Chakroun, M. Mezma, and D. Tuytens. “A GPU-accelerated Branch-and-Bound Algorithm for the Flow-Shop Scheduling Problem.” *CoRR*, vol. abs/1208.3933, 2012. [Online] Available at <http://arxiv.org/abs/1208.3933> [Accessed 17 July 2015].
- [21] I. Chakroun, M. Mezma, N. Melab, and A. Bendjoudi. “Reducing thread divergence in a GPU-accelerated branch-and-bound algorithm.” *Concurrency and Computation: Practice and Experience*, 2012. [Online] Available at <https://hal.inria.fr/hal-00731859> [Accessed 17 July 2015].
- [22] I. Chakroun and N. Melab. “Operator-level GPU-Accelerated Branch and Bound Algorithms.” *Procedia Computer Science*, vol. 18, pp. 280 – 289, 2013. 2013 International Conference on Computational Science.

## A Input Pre-processing Examples

This appendix provides an example of how the input is broken down into smaller problems which can then be solved by integer linear programming. The example proceeds as per the steps outlined in Algorithmalg:proc.

Table 1: An example input matrix, M, for the haplotype assembly problem. Each row of the matrix is a read,  $r_s$  is the start position of the read, and  $r_e$  is the end position of the read,  $h$  and  $h'$  are the haplotypes to be inferred from the input, and are shown for reference.

reads	1	2	3	4	5	6	7	$r_s$	$r_e$
$r_1$	0	-	-	-	-	-	-	1	1
$r_2$	1	0	-	-	-	-	-	1	2
$r_3$	0	0	0	-	-	-	-	1	3
$r_4$	0	1	-	-	-	-	-	1	2
$r_5$	-	-	0	1	-	-	-	3	4
$r_6$	-	-	1	0	-	-	-	3	4
$r_7$	-	-	-	-	-	1	1	6	7
$r_8$	-	-	-	-	-	0	0	6	7
$r_9$	-	-	-	0	1	-	0	4	7
$r_{10}$	-	-	-	-	0	0	0	5	7
$h$	1	0	0	1	0	0	0	-	-
$h'$	0	1	1	0	1	1	1	-	-

### A.1 Block Decomposition

The process of block decomposition is to split a large input matrix, say M, into smaller disjoint matrices, say  $C_k$  where k is the index of the submatrix. Formally, as per [5] block decomposition is computed as follows (reference will be made to the above input matrix).

Suppose that M is an input matrix with  $l$  columns (in this case 7), then for an integer  $j$ , whith  $1 < j < l$ , if there is no read,  $r$ , in M such that  $j$  is greater than the start position of  $r$  ( $r_s$  above) but less than the end position of  $r$  ( $r_e$  above), then column  $j$  of M (denoted  $M[j]$ ) is a splittable column. Suppose that there are  $k$  splittable columns in M, then the unsplittable blocks of M are formed by  $M[1, j_1]$ ,  $M[j_1, j_2]$ , ...,  $M[j_k, l]$ .

Consider the input matrix given in Table 1, all 7 columns of the input matrix must be checked to determine if they are splittable.

$j = 1$  For column 1 to be splittable, none of

the reads in the input must have  $r_s < 1 < r_e$ , which is true, hence  $M[1]$  is a splittable column.

$j = 2$  For column 2 to be splittable, none of the reads in the input must have  $r_s < 2 < r_e$ , however,  $r_3$  has a start position of 1 and an end position of 3, thus  $M[2]$  is not splittable.

The same process is applied for  $j = 3, 4, 5, 6$  and the splittable columns are found to be  $M[3]$  and  $M[4]$ . The paralellization of this step comes from performing the computation for all  $j$  simultaneously, thus the number of parallel processes would be equal to the number of columns in M, which would provide significant speed up when the input is large.

From the splittable columns (1, 3, 4), the unsplittable blocks are given by  $M[1, 1]$ ,  $M[1, 3]$ ,  $M[3, 4]$ ,  $M[4, 7]$ , and after removing irrelevant rows (rows containing only -'s) the unsplittable blocks are given by

$$\begin{aligned}
 M[1, 1] &= \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix} \\
 M[1, 3] &= \begin{bmatrix} 0 & - & - \\ 1 & 0 & - \\ 0 & 0 & 0 \\ 0 & 1 & - \\ - & - & 0 \\ - & - & 1 \end{bmatrix} \\
 M[3, 4] &= \begin{bmatrix} 0 & - \\ 0 & 1 \\ 1 & 0 \\ - & 0 \end{bmatrix} \\
 M[4, 7] &= \begin{bmatrix} 1 & - & - & - \\ 0 & - & - & - \\ - & - & 1 & 1 \\ - & - & 0 & 0 \\ 0 & 1 & - & 0 \\ - & 0 & 0 & 0 \end{bmatrix}
 \end{aligned}$$

### A.2 Singleton Row Removal

From the unsplittable blocks, singleton rows can be removed to reduce the problem further, since their removal does not modify the MEC score. A singleton row is a row for which the start and end position of a read are the same ( $r_s = r_e$ ). Informally this is a row for which there is only a single value. It follows that if an unsplittable block has only a single column, then all rows are singleton rows, as can be seen in  $M[1, 1]$  above.

These blocks do not need to be solved since they are included in other unsplittable blocks (M[1,1] is present in M[1,3]).

For M[1,3] the singleton rows are first, second to last, and last rows, their removal results in

$$M[1,3] = \begin{bmatrix} 1 & 0 & - \\ 0 & 0 & 0 \\ 0 & 1 & - \end{bmatrix}$$

For M[3,4] the singleton rows are the first and last, their removal results in

$$M[3,4] = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$$

For M[4,7], the singleton rows are the first and second, their removal results in

$$M[4,7] = \begin{bmatrix} - & - & 1 & 1 \\ - & - & 0 & 0 \\ 0 & - & 1 & 0 \\ - & 0 & 0 & 0 \end{bmatrix}$$

This offers two areas of parallelization, firstly across the unsplittable blocks. In this case there would be 3 parallel processes, one for M[1,3], another for M[3,4] and another for M[4,7], where each process would be eliminating the singleton rows within the block. Secondly, for each row within the unsplittable block. For process is created for each row in the block to determine if the row is a singleton, and to remove the row if it is a singleton.

### A.3 Duplicate Removal

For each row and column of an unsplittable block, the number of other rows or columns in the unsplittable block which are identical to it need to be determined. Furthermore, if an identical row or column is found, the rows or columns are merged and the multiplicity is increased. The multiplicity of the  $i^{th}$  row in an unsplittable block is denoted by  $w_i$ , the set of multiplicities for the rows by  $W$ , the multiplicity of the  $j^{th}$  column by  $c_j$ , and the set of multiplicities for the columns by  $C$ .

Using M[3,4] as an example, none of the rows or columns are duplicates, hence the multiplicities are

$$\begin{aligned} w_1 &= 1, w_2 = 1, W = \{1, 1\} \\ c_1 &= 1, c_2 = 1, C = \{1, 1\} \end{aligned}$$

This can also be parallelized by calculating the multiplicities for the columns and rows simultaneously.

## B Integer Linear Programming Example

This appendix demonstrates an example of how the haplotype assembly problem can be solved in parallel. The unsplittable block M[3,4] from Appendix A will be used.

Two sets of integers are defined which are used in the ILP example,  $J_{i,0}$  is the set of integers  $j \in \{1, 2, \dots, q\}$ , where  $q$  is the number of columns in the unsplittable block, such that the  $i^{th}$  entry in the  $j^{th}$  column of the unsplittable block is a 0, similarly  $J_{i,1}$  is the set of integers  $j \in \{1, 2, \dots, q\}$  such that the  $i^{th}$  entry in the  $j^{th}$  column of the unsplittable block is a 1.

For the unsplittable block M[3,4], the sets of integers would be

$$\begin{aligned} J_{1,0} &= \{1\} & J_{1,1} &= \{2\} \\ J_{2,0} &= \{2\} & J_{2,1} &= \{1\} \end{aligned}$$

## C Heuristic for Lower Bound Estimation