Parallelizing the Haplotying Assembly Problem

Author : Robert Clucas - 541590
Partner : Sasha Naidoo - 607993

July 26, 2015

Abstract

This report details the haplotype assembly problem, specifically the minimum error correction formulation. An algorithm for solving the minimum error correction formulation of the problem is chosen for parallelization which is based on integer linear programming and is known to be NP hard. The branch and bound algorithm is chosen as the method for solving the problem. Possible parallel implementations of both input data pre-processing and the branch and bound algorithm are provided, both of which aim to reduce the long run times by current implementations required to solve the problem. The project will be implemented using multi-core CPUs, Intel Xeon Phi coprocessors, and Nvidia GPUs. The software will be written using C and C++, and only free software will be used. The CUDA and OpenMP libraries will be used for the Nvidia GPUs and for communication between CPU threads respectively.

1 Introduction

It is commonly accepted that all humans share $\sim 99\%$ of the same DNA, however, small variations cause human beings to have different physical traits. Single nucleotide polymorphisms (SNPs), which are variations of a single DNA base from one individual to another, are believed to be able to address genetic differences. For diploid organisms, which have pairs of chromosomes, a *haplotype* is a sequence of SNPs in each copy of a pair of chromosomes. A *genotype* describes the conflated data of the haplotypes on a pair of chromosomes. Haplotypes are believed to contain more generic information than genotypes [1], however, obtaining haplotypes correctly is a difficult problem, which has been studied in two main forms : haplotype inference and haplotype assembly (HA).

Haplotype inference uses the genotype of a set of individuals. The genotype data tells the status of each allele at a position, but does not distinguish which copy of the chromosome the allele came from. The negative aspects of this approach are that it cannot distinguish rare and novel SNPs [2], and there is no way of knowing if the inferred haplotype is completely correct.

Haplotype assembly uses fragments of sequences generated by sequencing technology to determine haplotypes. The fragments of a sequence come from the two copies of an individual's chromosome. The goal of the haplotype assembly problem is then to correctly determine two haplotypes, where each haplotype corresponds to one of the two copies of the chromosome. Table 2 in Appendix A shows an example input, where h and h' are the assembled haplotypes.

The haplotype assembly problem was proven to be NP hard [3]. The algorithms used to solve the problem are thus computationally expensive and until recently, there was no practical exact algorithm to solve the minimum error correction (MEC) formulation of the problem [4], However, recently an exact algorithm was proposed by [5] which is capable of solving the MEC formulation exactly, and can thus correctly infer all haplotypes from the fragment sequences. Due the problem being NP hard, the implementation requires a long time to solve the problem - in the range of days for chromosomes with high errors rates. A parallel implementation could reduce the long run time, allowing useful haplotype information to be more quickly inferred from the available datasets, having positive effects in fields such as drug discovery, prediction of diseases, and variations in gene expressions, to name a few.

Parallel programming makes use of devices which have numerous cores, and uses these cores to execute a single instruction on multiple data (SIMD). The effectiveness of parallel programming is dependant

on the nature of the problem, as per Amdahl's law. While using multiple processors can potentially provide large performance increases, in practise the performance improvements are difficult to achieve due to additional complexities which are introduced by parallelism. The main difficulties are the synchronization and communication between the multiple cores, and the management of memory between the host (normally a CPU) and the device (parallel capable hardware, a GPU for example). Threads are created to operate on some data in parallel, and a block of threads is allocated memory which is accessible only to the threads in the block, while all threads from all blocks have access to a global memory space, however, access to this memory space is slower, decreasing performance. Furthermore, race conditions, where multiple threads attempt to access data at the same location simultaneously, can cause undefined behaviour and hence inaccurate results.

With the increasing popularity of General-Purpose GPU (GPGPU) programming, however, API's like CUDA [6] and OpenCL [7] which provide access to GPUs through simple function calls in C and C++, have made parallel programming easier on GPUs. Intel have also started to provide parallel focussed hardware, introducing the Xeon Phi, which, like a GPU, has numerous cores which can operate on data in parallel. The Xeon Phi is also programmed using an API [8], however, far more extensive use is made of compiler directives than by CUDA or OpenCL, with API functions being provided to gain access to parallel variables such as the thread index and number of threads which are running in parallel. Using these APIs, it is possible to write parallel implementations of algorithms which have significantly shortened run times.

The contribution of this paper is to consider the proposed algorithms for solving the HA problem, then to choose one for which parallelization can potentially provide performance improvements, and then to identify components of the selected algorithm which would be suited for parallelism, as well as proposing possible parallel implementation of the identified parallelizable components. Additionally, this paper details the specifications of the project and breakdown of work such that the project will be completed by the deadline.

The remainder of this paper is structured as follows. Section 2 provides background information related to the HA problem, some of the proposed algorithms for solving it, the MEC formulation, and the branch and bound algorithm. Section 3 describes the specifics of the project. Section 4 presents the general procedure for solving the chosen algorithm. Section 5 describes possible parallelization of the preprocessing of the input data. Section 6 describes the parallelization of the branch and bound algorithm, which is the chosen algorithm for solving the problem. Section 7 concludes.

2 Background

2.1 Haplotype Asembly Problem

This subsection will provide a brief overview of the haplotype assembly problem, and define the notation used through the rest of the paper. The input to the problem is a set of reads from a given genome sequence, where each read contains fragments from each of the two chromosomes which make up the genome sequence. The characters of a read consist of elements from a *ternary string*, where a ternary string has characters from the set {0, 1, -}. A value of 0 refers to the major allele at a site, a value of 1 to the minor allele, and a value of - to the lack of a read at the site, and is referred to as a *gap*. These reads are then combined to form a matrix, where each row of the matrix corresponds to a read (Table 2 in Appendix A provides an example).

Each column of the matrix is known as an SNP site. At each site, the data could be accurate, missing, or have error. The goal of the haplotype assembly problem is to determine a haplotype, H = {h,h'} from the matrix. The following terminology will be used to refer to properties of the input matrix and the fragments.

For the input matrix, M, the number of fragments is denoted by m, which is the number of rows in M. The number of SNP sites is denoted by n, which is the number of columns in M, while the j^{th} site of the i^{th} fragment is given by f_{ij}. Two fragments are said to conflict if the following conditions are true:

- $f_{ik} \neq f_{jk}$ **and** $f_{ik} \neq$ '-' **and** $f_{jk} \neq$ '-'

Essentially this means that for two fragments i and j, if at an SNP site k, the reads do not have error and are not gaps, the reads have different values (fragment i has a value 0 at site k, while fragment j has a value 1 at site k, or vice versa).

Following the notation of a conflict, the *distance* between two fragments is denoted by $d(f_i, f_j)$, and is

the total number of positions for which the two fragments f_i and f_j conflict. Furthermore, to understand some of the problem formulations from the fragment data, it is useful to define a *conflict graph* G = {V, E} [9], where V corresponds to a fragment, and E corresponds to an edge between two fragments if they conflict. If the input matrix contains no errors, then none of the fragments from the same chromosome will conflict and G will be bipartate. However, if there are errors (as is often the case) in M, G will not be bipartate. The haplotype assembly problem then requires the correction of G from a non-bipartate graph to a bipartate graph, from which two haplotypes can be assembled. There are numerous methods for solving the problem:

- **Minimum Fragment Removal (MFR):** This invloves removing the least number of fragments from the input data such that the resultant graph G is bipartate. It is shown in [9] that this can be solved in polynomial time.
- **Minimum Edge Removal (MER) [10]:** This method was recently proposed, and requires determining the minimum number of edges to remove such that removal of the edges results in G being bipartate.
- **Longest Haplotype Reconstruction (LHR):** This requires finding a set of fragments which, when they are removed from M, result in G being bipartate and the length of the resultant haplotypes being maximized [11].
- **Minimum Error Correction (MEC):** This method involves correcting the minimum number of elements (sites for all fragments) in M which allows the graph G to be bipartate. Although being the most complex method, it is the most widely used method as it provides the highest accuracy. Only recently has an exact algorithm for the MEC formulation been proposed which can provide an exact solution for all cases.

Due to the accuracy and more extensive previous work, the MEC formulation of the HA problem was chosen for a parallel implementation, hence will be the focus for the remainder of the paper.

2.2 Minimum Error Correction Formulations

Much research has been done on solving the MEC formulation of the HA problem. The first exact algorithm for solving the problem was a branch and bound algorithm proposed by [12]. The algorithm creates a tree which covers the search space of all possible corrections. The tree is then traversed to find the best solution. They use an upper bound which allows branches to be pruned when a better solution than the current best cannot result from further exploration of the branch, resulting in faster run times. However, this exact method has time complexity of $O(2^m)$, where m is the number of fragments in the input data, and hence cannot determine solutions for large problem sizes within a feasible amount of time. They also provide a heuristic method which uses a genetic algorithm to improve the computational time, with run times up to three orders of magnitude faster than the branch and bound implementation. While the heuristic method gives very similar results to the branch and bound implementation, it is slightly less accurate.

A dynamic programming solution was proposed by [13] which improved the run time problem for large input sizes. The algorithm has time complexity of $O(mk3^k + mlogm + mk)$, where k is the maximum number of SNP sites a fragment covers. In practice k is usually small, and results were shown small k (less than 100). The proposed dynamic programming solution had significant run time improvements over the solution proposed by [12]. However, for larger k values, the algorithm cannot solve the MEC formulation of the HA problem within a feasible amount of time.

More recently, [5] proposed an exact algorithm for solving the MEC problem. The proposed algorithm is the currently the only algorithm which can solve the HA problem for both the *all heterozygous* case (the alleles at an SNP site are assumed to be different) and the *general* case (the alleles at an SNP site may be the same due to error). Most other work assumes that the input fragment data is heterozygous, which, while true for most of the SNP sites in the input matrix, is often false for a small number of the SNP sites, which is the benefit of the algorithm proposed by [5]. Their exact algorithm removes unnecessary data from the input matrix, and partitions the matrix into smaller sub-matrices (or unsplittable blocks to use their terminology, Appendix A provides an example) which can then be solved in parallel. The problem is then formulated as an *integer linear programming* (ILP) problem and solved as a minimization problem. Their implementation was tested using an Intel i7-3960X CPU. For the HuRef dataset, the implementation required 12 days to solve the general case, and 31 hours for the all heterozygous case.

Heuristics methods are also proposed which speed up the computation by up to 15 times. However, the results are only shown for smaller input sizes, and the heuristic methods does not always determine the optimal solution. Most importantly, the solutions for the general case show lower MEC scores, meaning that the all heterozygous assumption is not always valid. This algorithm for the MEC formulation was chosen as it achieved optimal solutions for the all heterozygous and general cases. Their exact formulation of the problem as an ILP problem for the al lheterozygous case is now given.

2.2.1 All Heterozygous Algorithm by Chen

This subsection will summarize the all heterozygous ILP formulation of MEC HA problem as it is the algoirhm which was chosen for parallelization, and is referred to throughout the remainder of the paper.

The Hamming distance between two fragments i and j, d $= (f_i, f_j)$, is used by [5] to determine the MEC score of a solution H $= \{h, h'\}$, and is number of SNP sites at which the fragments conflict. Using the Hamming distance between all the fragments and the solution, the MEC score is given by

$$\text{MEC score} = \sum_{i=1}^{m} min\{d(f_i, h), d(f_i, h')\} \quad (1)$$

Where m is again the number of rows in the input matrix M. The MEC score is *optimal* if it is the minimum possible score. Some assumptions are made for the input matrix as per [5], namely

- No row of the input matrix is useless - at least one entry in the row is a 1 or 0
- No column of the input matrix is monotone - the column must have at least one 0 and one 1
- No column of M contains more 1's than 0's - if this is not the case the values are all flipped, which does not change the solution or the MEC score

The input matrix first undergoes pre-processing to reduce the size of the input and to break the input into multiple blocks which can be solved in parallel. It must be noted that the pre-processing does not affect the solution in any way. The pre-processing functions which are applied are:

- **Block decomposition:** This process takes the input matrix and splits it into smaller, unsplittable blocks, which are disjoint and can be solved independently.

- **Singleton removal:** A singleton is a row for which the start and end positions of the fragment are the same - i,e the fragment has only one element. Since singletons do not affect the MEC score, they can be removed. This is done for all rows in all the smaller, unsplittable blocks.
- **Duplicate removal:** Rows and columns which are the same are merged into a single row or column, and the multiplicity of the row or column is recorded. This is done for all rows and all columns in all the smaller, unsplittable blocks.

Additionally for the general case, the intrinsically heterozygous columns need to be determined. Determining if a column j is intrinsically heterozygous requires counting the number of rows in the input matrix which have a value of 0 (1) at position j, which is stored as a variable $n_{j,0}$ ($n_{j,1}$), as well as the number of non-singular rows in the input matrix which have a value of 0 (1) at position j, which is stored as a variable $s_{j,0}$ ($s_{j,1}$). Column j is then intrinsically heterozygous if

$$min\{n_{j,0}, n_{j,1}\} \geq \left\lceil \frac{s_{j,0} + s_{j,1}}{2} \right\rceil$$

From the reduced blocks the ILP problem is formulated for the all heterozygous and the general case. The all heterozygous case formulation will be shown, the general case formulation is shown in [14].

2.2.1.1 ILP Formulation For integers $p, q \in \mathbb{Z}$, where $1 \leq p \leq m$, and $1 \leq q \leq n$, p is the index of the fragment (or row) in the input matrix, M, and q is the index of SNP (or column) in M. The multiplicity of the j^{th} column of M is denoted by c_j, while the multiplicity of the i^{th} row of M is denoted by c_j. The following binary variables are introduced, y_i, which has a value of 1 if $d(f_i, \text{h}) \leq d(f_i, \text{h'})$, otherwise has a value of 0 (when $d(f_i, \text{h'}) < d(f_i, \text{h})$). Informally, if fragment i is part of h then y_i has a value of 1, otherwise it has a value of 0. x_j, which has a value of 1 if the j^{th} bit of h is 1, otherwise has a value of 0 (if the j^{th} bit of h is 0). Lastly $J_{i,0}$ ($J_{i,1}$) are the sets of integers $j \in \{1, 2, ..., q\}$ for which the i^{th} value in column j is a 0 (1). An example of the ILP formulation is is given in Appendix B.

Using the above defined variables, an integer programming formulation is possible, however, a non-linear term in the form of $y_i x_j$ arises, which cannot be solved using ILP techniques. To overcome this problem, [5] defines a variable $t_{i,j}$ for $y_i x_j$ and im-

poses constraints on the variable which ensure linearity (the constraints are shown in the final formulation of the problem below). The final ILP formulation of the HA problem for the all heterozygous case is

$$\text{Minimize} \quad \sum_{i=1}^{p} w_i \sum_{j \in J_{i,0}} c_j(1 - x_j - y_i + 2t_{i,j})$$
$$+ \sum_{i=1}^{p} w_i \sum_{j \in J_{i,1}} c_j(y_i + x_j - 2t_{i,j})$$

$$\text{Subject to} \quad \forall_{1 \le i \le p} \;\; y_i \in \{0,1\}$$
$$\forall_{1 \le j \le q} \;\; x_j \in \{0,1\}$$
$$\forall_{1 \le i \le p} \;\; \forall_{1 \le j \le p} \;\; t_{i,j} \in \{0,1\}$$
$$t_{i,j} \le y_i$$
$$t_{i,j} \le x_j$$
$$t_{i,j} \ge y_i + x_j + 1$$

2.3 Branch and Bound

The branch and bound algorithm divides the search space into sub spaces using a branching operator, then each sub space is explored for the optimal solution. Each of these sub spaces is represented as a branch on a tree, a bounding operator is used to determine the best possible solution the branch can provide. A pruning operator is used to remove branches which cannot provide a solution which is better than the current best, based on the result of the bounding operator. The branch and bound algorithm can be parallelized in multiple ways [15]. Either specific, computationally difficult functions can be accelerated by a parallel implementation (for example the lower bound calculation), or the entire tree can be searched in parallel, or a combination of both can be employed. Using tree based approaches will require communication between the processes searching a branch, as the solution of each branch will need to be compared to the solutions of the other branches to determine the globally optimal solution. This can become a bottleneck for both CPU and GPU implementations if the work is not divided correctly, as processes may have to wait for other processes. There have been numerous parallel branch and bound algorithms proposed for both the CPU and GPU which deal with these problems in different ways.

2.3.1 CPU Implementations

The ALPS framework [16] is written in C++ and provides a parallel CPU implementation of the branch and bound algorithm. The framework is tested on the knapsack problem [17], which is an ILP problem and is known to be NP hard. The speedup achieved is near linear in the number of nodes when the number of nodes is small, but diminishes as more nodes are added due to the amount of communication required between the nodes. A speedup of 8 times is achieved for 8 nodes, and 26 times for 32 nodes. It is likely that for the HA problem the number of nodes could be extremely large, in which case this method will be infeasible.

The MALLBA framework [18] is also written in C++ and provides a parallel branch and bound algorithm. It labels nodes as *master* or *slave* nodes, which defines the type of work the nodes do. The framework also provides heuristic methods for solving the problems more efficiently, however, the accuracy of the results is lowered, which while acceptable for many applications, is not acceptable for the HA problem. The performance results are similar to the ALPS framework, where near linear speed up is achieved for a small number of nodes, but levels off for larger numbers of nodes.

2.3.2 GPU Implementations

A heterogeneous CPU-GPU implementation was proposed by [19] and was applied to the knapsack problem. Due to the overhead of transferring data from the CPU to the GPU before computation, the model only uses GPUs when the tree has a large number of nodes (> 5000), otherwise the CPUs are used. The tree is built using a *breadth first* strategy to favour the parallel nature of the GPUs. Both the branching and bounding steps can be performed on the CPU or GPU. If the number of nodes is sufficiently large such that full occupancy of the GPU is ensured, then for each iteration of the algorithm the GPU does the branching and bounding on a list of nodes, eliminates nodes for which an optimal solution cannot be found, and returns the list back to the CPU. This process continues until convergence. This regular communication between the CPU and GPU is expensive, which lessens the speedup achieved by the implementation. However, a speedup of up to 9.27 times was achieved for larger problem sizes, validating the feasibility of the branch and bound algorithm for parallel implementation. It must be noted that the algorithm was developed for Nvidia's Fermi architecture, which does not support dynamic parallelism [20], thus requires numerous CPU-GPU communication. The newer Kepler architecture, however, does

support dynamic parallelism and could minimise this communication.

In [21, 22], algorithms are proposed which target the bounding operation for parallelism, as well as dealing with the problem of thread divergence when using GPUs for branch and bound, which comes from the position of the nodes in the tree, and the need for branches to communicate with other branches to compare solutions. Depending on the problem to which branch and bound is applied, the tree structure can be irregular, which makes paralellizing the tree search difficult for GPUs not supporting recursion, as was the case when the algorithm was proposed – hence the focus on only the bounding operation. The algorithm was applied to the Flow-Shop scheduling problem, for which it is shown that ∼98% of the computation is spent on the bounding operation. Speed ups of up to 100 times were achieved by the GPU implementation over a multi-threaded CPU implementation. However, this kind of speedup will only be seen in problems where the calculation of the bounds is the bottleneck. Nevertheless, even if significantly less time is spent calculating the bounds, the speed up should still be considerable.

The algorithms of [21, 22] are extended by [23] to include not only parallelization of the bounding operator, but also the branching and pruning operators. This implementation uses the GPU for almost all the searching of the subspaces. However, the CPU is still used for the comparison of the solutions between each iteration as each GPU thread can only determine the solution of its first child nodes, rather than all child nodes until the leaves of the tree. Again this is due to the limitations of the Fermi architecture not supporting recursion. Despite this hardware limitation, the algorithm is still able to achieve a speed up of up to 166 times over a multi-threaded CPU implementation, for large problem sizes.

3 Project Description

Based on the related work reviewed in Section 2, the specification of the project and its purpose, the project schedule, and the assumptions, constraints and requirements, are given in the subsections to follow.

3.1 Project Specification

The aim of the project is to implement the HA problem in parallel on 3 different hardware configurations

Table 1: Project schedule detailing milestones and the dates by which they should be completed.

Date	Milestone
3 Aug	Get access to required hardware
10 Aug	Install all required software
4 Sept	CPU OpenMP/MPI implementation
7 Oct	GPU and Xeon Phi Implementation
15 Oct	Open day - implementations debugged
23 Oct	Report and presentation complete

and then to compare the performance of the configurations. The chosen algorithm for parallel implementation is the MEC formulation proposed by [5] since the algorithm can solve the all heterozygous and general case of the HA problem optimally. The 3 configurations are:

- CPU based cluster
- Intel Xeon Phi Coprocessor based cluster
- Nvidia GPU based cluster

3.2 Project Schedule

The schedule presented in Table 1 details the milestones which have been set to ensure that the project is completed on time.

3.3 Hardware

The following hardware will be required for the project:

- A cluster with multi-core CPUs
- A custer with at least one Intel Xeon Phi Coprocessor
- A cluster with at least one Nvidia GPU with compute capability ≥ 3.5

3.4 Software

The following software will be used for the project:

- C and C++, as these are the languages used for programming Intel Xeon Phi's and Nvidia GPU's
- The Nvidia CUDA API for programming the Nvidia GPUs
- The OpenMP API for the CPU and Intel Xeon Phi implementations
- The Nvidia nvcc compiler for compiling the GPU code to run on the Nvidia GPU

- The Intel compiler suite for compiling the CPU and Intel Xeon Phi implementations

3.5 Assumptions

The following assumptions are made for the project:

- The hardware detailed in Section 3.3 is available
- The software detailed in Section 3.4 is available
- The datasets for the SNP inputs will all be available
- Solutions for the datasets are available to verify the corectness of the implementations

3.6 Constraints

The following constraints are placed on the project:

- The available budget is limited and thus a limited number of hardware devices will be available
- All software to be used must be free
- The project must be completed by Robert Clucas and Sasha Naidoo

3.7 Success Criteria

Due to the difficulty of the problem, the following criteria are defined which must be met for the project to be considered a success:

- All three implementations are working by the project demonstration date
- All implementations can at least solve the all heterozygous case correctly, with optimal MEC scores
- At least one implementation provides a speed up over the results detailed in [5]

4 Procedure for Solving MEC HA

Using the information provided by Section 2.2.1 and Section 2.2.1.1, the general procedure for solving the MEC HA problem proposed by [5] is detailed by Algorithm 1.

5 Parallel Pre-processing for MEC HA

Observing the procedure for solving the MEC HA problem presented by Algorithm 1, there are numerous areas to which parallelism can be applied.

Algorithm 1 Procedure for solving the MEC HA problem using ILP

Step 1: Perform block decomposition on columns in M to obtain smaller unsplittable blocks C_k (if general case, check if column is intrisically heterozygous)

Step 2: Remove singleton rows from all C_k unsplittable blocks (if general case, check if singleton row starts and ends on intrinsically heterozygous columns)

Step 3: Remove duplicate rows and columns from all C_k unsplittable blocks

Step 4: Solve ILP formulation using branch and bound algorithm on all C_k unpsplittable blocks

Step 5: Concatenate solutions of all C_k unsplittable blocks

5.1 Determining Intrinsically Hetrozygous Columns

Determining intrinsically heterozygous columns in M requires computing, for all j columns in M, the variables $n_{j,0}$, $n_{j,1}$, $s_{j,0}$, $s_{j,1}$ as described in Section 2.2.1. A parallel implementation could use two threads for each column, where one thread would compute the values of the n variables, while the other would compute the values of the s variables. This would require $2n$ threads, where n is the number of columns in M.

5.2 Singleton Removal

Finding singletons is a simple task as it only requires determining if there is more than one element in a row. Each thread checks if its corresponding entry in the row is a 0 or 1 and, if it is, increments a counter. If the value of the counter is > 1, then the row is not singular, since more than one element in the row must have had a value of 0 or 1. This would have time complexity of O(1) on the GPU, compared to O(m) using the implementation of [5]. Additionally, vectorized operations could be used for further performance increase if the data is stored contiguously.

5.3 Duplicate Removal

For the implementation used by [5], each row needs to be compared with each other row to determine whether the rows are identical, and the same procedure is required for the columns. This can be done with time complexity of O(Llogk), where L is the total length of the reads and k is the number of reads. If the total length of the reads is large, this process can require large amounts of time for computation. The parallel approach for this operation would have

two components.

Simultaneous column and row search: This would compute the duplicate rows and columns simultaneously.

Parallelism within row and column search: Multiple rows and columns can be compared at the same time. Consider N rows or columns, then $\frac{N}{2}$ threads could be used where the i^{th} thread checks if its row or column is identical to the $(i + \frac{N}{2})^{th}$ row or column. This would continue iteratively, where the number of threads would half on each iteration until completion. The time complexity would be logarithmic for the column and row search, with the total time complexity being $\max\{log\text{m},log\text{n}\}$, where m and n are the number of rows and columns to check for duplicates, respectively.

6 Parallel Branch and Bound for MEC HA

Considering the related work presented in Section 2.3, a similar approach to those presented in [21–23] will be taken, however, an attempt will be made to limit the number of transactions between the CPU and GPU, preferring rather to perform as much of the computation on the GPU as possible. Furthermore, the implementation will not attempt to balance the tree since the parallel devices which will be used are capable of recursion, allowing branches to be explored until either an optimal or infeasible solution is reached.

6.1 Branching Operator

The branching operator is used to create child nodes from the current node, and continues recursively until the leaves of the branch are found. The operator determines which of the variables should be branched on, and the selection operator is called to request additional threads so that the child nodes can be explored in parallel.

6.2 Selection Operator

Since the algorithm proceeds in parallel, the selection operator is not used to decide which branch must be explored next, but rather to decide which branch should be allocated more threads first when multiple branches request additional threads. The selection operator uses the lower bound of the branch to determine which branch is given additional threads

first, hence the selection strategy is that of best first. The selection operator also keeps track of the total number of threads in use so that threads are only allocated if they are available.

6.3 Pruning Operator

The pruning operator removes branches which have a lower bound greater than the current best solution, and is applied to each of the branches, which allow it to operate in parallel. Additionally the pruning operator moves the threads of the pruned branches back into the thread pool so that they can be allocated to other branches by the selection operator.

6.4 Bounding Operator

The bounding operator is the most important component of the parallel branch and bound implementation. It is used to determine the lower bound of a branch, which gives an estimate for the best solution a branch can provide. The bounding operator involves computing the current score of the branch using the values of the known variables. Since adding more variables can only increase the score of the branch, this provides a lower bound for the score. A limitation, however, is that the lower bound is only relevant when compared to another lower bound computed using at least the same variables. For example if the bound is computed using x_1, y_1, and y_2, then the bound can only be compared to another bound which used either the same variables, or the same variables as well as additional variables. This will always be the case when a branch is traversed downwards, as is the case for a recursive implementation, and is therefore not a problem. The current score can be computed in parallel, since the minimization problem described in Section 2.2.1.1, involving two sets of multiple summations, can be created for each of the outer sums (the sums with the index i in the ILP formulation).

6.4.1 Current score

If, for the thread index i the corresponding variable y_i is known, then a thread creates two more threads, where one thread computes

$$\sum_{j \in J_{i,0}} c_j(1 - x_j - y_i + 2t_{i,j}) \quad \forall \text{ known } x_j$$

and the other computes

$$\sum_{j \in J_{i,1}} c_j(y_i + x_j - 2t_{i,j}) \quad \forall \text{ known } x_j$$

The constraint for all known x variables is required, since both x_j and y_i must be known to determine the score. The variable $t_{i,j}$ can be determined from the constraints if both x_j and y_i are known. The results of each thread are then added to give the current score over the know variables for the thread index.

Once all the threads are complete, the score of each of the threads is added to give the total current score, which is the lower bound of the branch.

7 Conclusion

Research into the haplotype assembly problem, specifically the minimum error correction formulation, as well as the branch and bound algorithm and its parallel implementations, has been conducted in this report. From this research, an integer linear programming formulation of the problem was chosen for parallelization. Areas of parallelization are identified for both the pre-processing of the input data and the branch and bound algorithm itself. The proposed parallelizations will theoretically reduce the long run times required by current implementations for solving the problem. There will be three implementations, one using multi-core CPUs, another using the Intel Xeon Phi coprocessor, and the last using Nvidia GPUs. The required hardware and software is outlined, and the project schedule is given.

References

[1] J. Stephens, "Haplotype variation and linkage disequilibrium in 313 human genes," *Science*, vol. 293, pp. 489–493, July 2001. [Online]. Available at `http://dx.doi.org/10.1126/science.1059431` [Accessed 27 June 2015].

[2] D. He, A. Choi, K. Pipatsrisawat, A. Darwiche, and E. Eskin, "Optimal algorithms for haplotype assembly from whole-genome sequence data," vol. 26, no. 12, pp. i183–i190, 2010.

[3] R. Lippert, R. Schwartz, G. Lancia, and S. Istrail, "Algorithmic strategies for the single nucleotide polymorphism haplotype assembly problem," vol. 3, no. 1, pp. 23–31, 2002.

[4] P. Bonizzoni, G. Della Vedova, R. Dondi, and J. Li, "The haplotyping problem: An overview of computational models and solutions," *Journal of Computer Science and Technology*, vol. 18, no. 6, pp. 675–688, 2003. [Online]. Available at `http://dx.doi.org/10.1007/BF02945456` [Accessed 27 June 2015].

[5] Z.-Z. Chen, F. Deng, and L. Wang, "Exact algorithms for haplotype assembly from whole-genome sequence data," vol. 29, no. 16, pp. 1938–1945, 2013.

[6] Nvidia, *CUDA C PROGRAMMING GUIDE. Version 7.0.* 2015. [Online] Available at `http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#axzz3gRwjcM3m` [Accessed 20 July 2015].

[7] Khronos, *The OpenCL Specification. Version 2.1.* 2015. [Online] Available at `https://www.khronos.org/registry/cl/` [Accessed 20 July 2015].

[8] Intel, *Intel Xeon Phi Coprocessor DEVELOPER's QUICKSTART GUIDE. Version 1.8.* 2013.

[9] G. Lancia, V. Bafna, S. Istrail, R. Lippert, and R. Schwartz, "Snps problems, complexity, and algorithms," in *Algorithms ESA 2001* (F. auf der Heide, ed.), vol. 2161 of *Lecture Notes in Computer Science*, pp. 182–193, Springer Berlin Heidelberg, 2001. [Online] Available at `http://dx.doi.org/10.1007/3-540-44676-1_15` [Accessed 16 July 2015].

[10] D. Aguiar and S. Istrail, "Hapcompass: A fast cycle basis algorithm for accurate haplotype assembly of sequence data," *Journal of Computational Biology*, vol. 19, pp. 577–590, June 2012. [Online] Available at `http://www.brown.edu/Research/Istrail_Lab/papers/hapcompass_jcb_draft.pdf` [Accessed 16 July 2015].

[11] R. Schwartz, "Theory and algorithms for the haplotype assembly problem," *Commun. Inf. Syst.*, vol. 10, no. 1, pp. 23–38, 2010. [Online] Available at `http://projecteuclid.org/euclid.cis/1268143371` [Accessed 16 July 2015].

[12] R.-S. Wang, L.-Y. Wu, Z.-P. Li, and X.-S. Zhang, "Haplotype reconstruction from snp fragments by minimum error correction," vol. 21, no. 10, pp. 2456–2462, 2005.

[13] M. Xie, J. Wang, and J. Chen, "A practical exact algorithm for the individual haplotyping problem mec," in *BioMedical Engineering and Informatics, 2008. BMEI 2008. International Conference on*, vol. 1, pp. 72–76, May 2008.

[14] Z.-Z. Chen, F. Deng, and L. Wang, "Suppementary material for exact algorithms for haplotype assembly from whole-genome sequence data," 2013. [Online] Available at `http://bioinformatics.oxfordjournals.org/content/suppl/2013/06/18/btt349.DC1` [Accessed 20 April 2015].

[15] T. G. Crainic, B. Le Cun, and C. Roucairol, *Parallel Branch-and-Bound Algorithms*, pp. 1–28. John Wiley & Sons, Inc., 2006. [Online] Available at `http://dx.doi.org/10.1002/9780470053928.ch1` [Accessed 17 July 2015].

[16] Y. Xu, T. Ralphs, L. Ladnyi, and M. Saltzman, "Alps: A framework for implementing parallel tree search algorithms," in *The Next Wave in Computing, Optimization, and Decision Technologies* (B. Golden, S. Raghavan, and E. Wasil, eds.), vol. 29 of *Operations Research/Computer Science Interfaces Series*, pp. 319–334, Springer US, 2005. [Online] Available at `http://dx.doi.org/10.1007/0-387-23529-9_21` [Accessed 17 July 2015].

[17] M. Kedia, "Dynamic programming," October 2005. [Online] Available at `http://www.cs.cmu.edu/afs/cs/academic/class/15854-f05/www/scribe/lec10.pdf` [Accessed 27 June 2015].

[18] E. Alba, F. Almeida, M. Blesa, J. Cabeza, C. Cotta, M. Daz, I. Dorta, J. Gabarr, C. Len, J. Luna, L. Moreno, C. Pablos, J. Petit, A. Rojas, and F. Xhafa in *Euro-Par 2002 Parallel Processing* (B. Monien and R. Feldmann, eds.), vol. 2400 of *Lecture Notes in Computer Science*, pp. 927–932, Springer Berlin Heidelberg, 2002. [Online] Available at `http://dx.doi.org/10.1007/3-540-45706-2_132` [Accessed 17 July 2015].

[19] A. Boukedjar, M. Lalami, and D. El-Baz, "Parallel branch and bound on a cpu-gpu system," in *Parallel, Distributed and Network-Based Processing (PDP), 2012 20th Euromicro International Conference on*, pp. 392–398, Feb.

[20] Nvidia, "Cuda dynamic parallelism programming guide," August 2012. [Online] Available at `https://www.clear.rice.edu/comp422/resources/cuda/html/cuda-dynamic-parallelism/index.html` [Accessed 17 July 2015].

[21] N. Melab, I. Chakroun, M. Mezmaz, and D. Tuyttens, "A gpu-accelerated branch-and-bound algorithm for the flow-shop scheduling problem," *CoRR*, vol. abs/1208.3933, 2012. [Online] Available at `http://arxiv.org/abs/1208.3933` [Accessed 17 July 2015].

[22] I. Chakroun, M. Mezmaz, N. Melab, and A. Bendjoudi, "Reducing thread divergence in a GPU-accelerated branch-and-bound algorithm," *Concurrency and Computation: Practice and Experience*, 2012. [Online] Available at `https://hal.inria.fr/hal-00731859` [Accessed 17 July 2015].

[23] I. Chakroun and N. Melab, "Operator-level gpu-accelerated branch and bound algorithms," *Procedia Computer Science*, vol. 18, pp. 280 – 289, 2013. 2013 International Conference on Computational Science.

The following two appendices are examples of the proposed work and are formally being declared as joint work.

A   Input Pre-processing Example

This appendix provides an example of how the input is broken down into smaller problems which can then be solved more easily as an integer linear programming problem. The example proceeds as per the steps outlined in Algorithm 1. The examples are shown for the all heterozygous case, for the general case, the block decomposition and singleton removal steps would require determining the intrinsically heterozygous columns.

Table 2: An example input matrix, M, for the haplotype assembly problem. Each row of the matrix is a read, r_s is the start position of the read, and r_e is the end position of the read, h and h' are the haplotypes to be assembled from the input, and are shown for reference.

reads	1	2	3	4	5	6	7	r_s	r_e
r_1	0	-	-	-	-	-	-	1	1
r_2	1	0	-	-	-	-	-	1	2
r_3	0	0	0	-	-	-	-	1	3
r_4	0	1	-	-	-	-	-	1	2
r_5	-	-	0	1	-	-	-	3	4
r_6	-	-	1	0	-	-	-	3	4
r_7	-	-	-	-	-	1	1	6	7
r_8	-	-	-	-	-	0	0	6	7
r_9	-	-	-	0	1	-	0	4	7
r_{10}	-	-	-	-	0	0	0	5	7
h	1	0	0	1	0	0	0	-	-
h'	0	1	1	0	1	1	1	-	-

A.1   Block Decomposition

The process of block decomposition is to split a large input matrix, say M, into smaller disjoint matrices, say C_k, where k is the index of the submatrix. Formally, as per [1] block decomposition is computed as follows (reference will be made to the above input matrix).

Suppose that M is an input matrix with l columns (in this case 7), then for an integer j, with $1 < j < l$, if there is no read, r, in M such that j is greater than the start position of r (r_s above) but less than the end position of r (r_e above), then column j of M (denoted M[j]) is a splittable column. Suppose that there are k splittable columns in M, then the unsplittable blocks of M are formed by M[1, j_1], M[j_1, j_2],..., M[j_k, l].

Consider the input matrix given in Table 2, all 7 columns of the input matrix must be checked to determine if they are splittable.

$j = 1$: For column 1 to be splittable, none of the reads in the input must have $r_s < 1 < r_e$, which is true, hence M[1] is a splittable column.

$j = 2$: For column 2 to be splittable, none of the reads in the input must have $r_s < 2 < r_e$, however, r_3 has a start position of 1 and an end position of 3, thus M[2] is not splittable.

The same process is applied for $j = 3$, 4, 5, 6 and the additional splittable columns are found to be M[3] and M[4]. The parallelization of this step comes from performing the computation for all j simultaneously, thus the number of parallel processes would be equal to the number of columns in M, which would provide significant speed up when the input in large.

From the splittable columns (1, 3, 4), the unsplittable blocks are then given by M[1, 1], M[1, 3], M[3, 4], M[4, 7], and after removing irrelevant rows (rows containing only -'s) the unsplittable blocks are

$$M[1,1] = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

$$M[1,3] = \begin{bmatrix} 0 & - & - \\ 1 & 0 & - \\ 0 & 0 & 0 \\ 0 & 1 & - \\ - & - & 0 \\ - & - & 1 \end{bmatrix}$$

$$M[3,4] = \begin{bmatrix} 0 & - \\ 0 & 1 \\ 1 & 0 \\ - & 0 \end{bmatrix}$$

$$M[4,7] = \begin{bmatrix} 1 & - & - & - \\ 0 & - & - & - \\ - & - & 1 & 1 \\ - & - & 0 & 0 \\ 0 & 1 & - & 0 \\ - & 0 & 0 & 0 \end{bmatrix}$$

A.2 Singleton Row Removal

From the unsplittable blocks, singleton rows can be removed to reduce the probelm furhter, since their removal does not modify the MEC score. A singleton row is a row for which the start and end position of a read are the same ($r_s = r_e$). Informally this is a row for which there is only a single value. It follows that if an unsplittable block has only a single column, then all rows are singleton rows, as can be seen in M[1,1] above. These blocks do not need to be solved since they are included in other unsplittable blocks (M[1,1] is present in M[1,3]).

For M[1,3] the singleton rows are first, second to last, and last rows, their removal results in

$$M[1,3] = \begin{bmatrix} 1 & 0 & - \\ 0 & 0 & 0 \\ 0 & 1 & - \end{bmatrix}$$

For M[3,4] the singleton rows are the first and last, their removal results in

$$M[3,4] = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$$

For M[4, 7], the singleton rows are the first and second, their removal results in

$$M[4,7] = \begin{bmatrix} - & - & 1 & 1 \\ - & - & 0 & 0 \\ 0 & - & 1 & 0 \\ - & 0 & 0 & 0 \end{bmatrix}$$

This offers two areas of parallelization, firstly across the unsplittable blocks. In this case there would be 3 parallel processes, one for M[1,3], another for M[3,4] and another for M[4,7], where each process would be eliminating the singleton rows within the block. Secondly, for each row within the unplittable block. A process is created for each row in the block to determine if the row is a singleton, and to remove the row if it is a singleton.

A.3 Duplicate Removal

For each row and column of an unsplittable block, the number of other rows or columns in the unsplittable block which are identical to it need to be determined. Furthermore, if an identical row or column in found, the rows or columns are merged and the multiplicity is increased. The multiplicity of the i^{th} row in an unsplittable block is denoted by w_i, the set of multiplicities for the rows by W, the multiplicity of the j^{th} column by c_j, and the set of multiplicities for the columns by C.

Using M[3,4] as an example, none of the rows or columns are duplicates, hence the multpilicities are

$$w_1 = 1, \; w_2 = 1, \; W = \{1, 1\}$$
$$c_1 = 1, \; c_2 = 1, \; C = \{1, 1\}$$

This can also be parallelized by calculating the multiplicities for the columns and rows simultaneously.

B Integer Linear Programming Example

This appendix demostrates an example of how the haplotype assembly problem can be solved in parallel, as well as how the branching and bounding operators are used. The unsplittable block M[3,4] from Appendix A will be used.

Two sets of integers are defined which are used in the ILP example, $J_{i,0}$ is the set of integers $j \in \{1,2,...,q\}$, where q is the number of columns in the unsplittable block, such that the i^{th} entry in the j^{th} column of the unsplittable block is a 0, similarly $J_{i,1}$ is the set of integers $j \in \{1,2,...,q\}$ such the the i^{th} entry in the j^{th} column of the unsplittable block is a 1.

For the unsplittable block M[3,4], the sets of integers would be

$$J_{1,0} = \{1\} \quad J_{1,1} = \{2\}$$
$$J_{2,0} = \{2\} \quad J_{2,1} = \{1\}$$

The variables to be solved are then x_1, x_2, y_1, and y_2. as per the ILP formulation described in Sec-


Figure 1: Tree illustrating the known variables for the nodes and the lower bounds determined using the variables. A white node represents a branch which should be explored, while a black node represents a branch which should be pruned.

z = [0 1 0 -]   ③   z = [0 1 1 -]
LB = 1   LB = 0

④   ⑤

Figure 2: Tree illustrating the known variables for the nodes and the lower bounds determined using the variables. A white node represents a branch which should be explored, while a black node represents a branch which should be pruned.

tion 2.2.1 in the main text. For the purposes of illustration, the set of variables will be represented by $\mathbf{z} = [x_1, y_1, x_2, y_2]$. The tree is first created with two roots by branching the first variable, x_1. Each root is created by a thread. One thread sets $x_1 = 0$, and the other sets $x_1 = 1$. Each of these threads then branches the next variable y_1, again with two threads. The rest of the example details the procedure for one of the branches, with $x_1 = 0$.

Figure 1 shows the root created with $x_1 = 0$, and then the branching of y_1. The calculation of the lower bound for the two child nodes is the done is parallel. Additionally the computation of each of the bounds is parallelized as explained in Section 6.4 of the main text. Two threads are used, one for the top summation over j, and another for the bottom summation over j. Thus for $i = 1$, the first thread uses $j = J_{1,0} = 1$ and computes

$$\text{partial sum} = c_1(1 - x_1 - y_1 + 2t_{1,1})$$
$$= 1(1 - 0 - 0) = 1$$

simultaneously, the other thread uses $j = J_{1,1} = 2$ and computes

$$\text{partial sum} = c_1(y_1 + x_2 + 2t_{1,2})$$
$$= 1(1 - ? - ?) = -$$

z = [0 1 1 0]   ⑤   z = [0 1 1 1]
S = 0   S = 2

⑥   ⑦

Figure 3: Tree illustrating the known variables for the nodes and the scores of each of the branches.

but since x_2 is not know, it does not return a value. The lower bound is then 1, since only the result of the first thread is used.

At the same time, the lower bound of the right branch would have been computed using the same procedure to determine the lower bound of 0 as shown in Figure 1. Since the lower bound of the left branch is higher than that of the right branch, the left branch is pruned and will not be explored further. Node 3, however, must be explored so 2 new threads are created for branching, and the same procedure is applied in parallel. The known variables and the resulting lower bounds for this are shown in Figure 2

Node 4 in Figure 2 has a lower bound of 1, which is higher than the lower bound of node 5, which is 0. Node 4 is pruned while new threads are created to explore node 5 further. Since there is only one variable left for branching, the two child nodes from node 5 return the score of the branch rather than the lower bound since the child nodes are leaves. Figure 3 shows the results of branching from node 5, and the scores of the branch.

Since the leaves of the branch have been found, the lowest score (in this case the score of node 6) is return to the root of the tree, where is is compared to the solutions of all the other branches to determine which solution is optimal. For this branch, the solution would be $x_1 = 0$, $x_2 = 1$, $y_1 = 1$, $y_2 = 0$, which is the optimal solution for all branches for this example. Since the concatination of the x variables is the solution for the haplotype h, $h_{3,4} = [x_1, x_2] = [0,1]$ is the assmebled component of the haplotype for the unsplittable block M[3,4], which can be verified by Table 2.

References

[1] Z.-Z. Chen, F. Deng, and L. Wang. Exact algorithms for haplotype assembly from whole-genome sequence data. 29(16):19381945, 2013.